

Data analytics (ACS0603)

UNIT-3 NOTES

Data preprocessing is a crucial step in data analytics, ensuring that raw data is cleaned, transformed, and structured for effective analysis. The main forms of data preprocessing include:

1. Data Cleaning

- Handling missing values (e.g., imputation, deletion)
- Removing duplicates
- Correcting errors and inconsistencies
- Handling outliers (e.g., using Z-score, IQR)

2. Data Integration

- Combining data from multiple sources
- Handling schema mismatches and entity resolution
- Merging datasets for a unified analysis

3. Data Transformation

- **Normalization:** Scaling values between a fixed range (e.g., Min-Max Scaling, Z-score normalization)
- **Standardization:** Converting data to have a mean of 0 and a standard deviation of 1
- **Encoding categorical data:** One-hot encoding, Label encoding
- **Feature engineering:** Creating new meaningful features from existing ones

4. Data Reduction

- **Dimensionality reduction:** PCA, t-SNE, LDA
- **Feature selection:** Selecting relevant features to reduce complexity
- **Sampling:** Random sampling, Stratified sampling

5. Data Discretization and Binning

- Converting continuous variables into categorical bins (e.g., age groups: young, middle-aged, senior)

- Helps in handling skewed distributions

Data Cleaning Techniques in Data Analytics

Data cleaning is a crucial step in data preprocessing to ensure that the dataset is accurate, consistent, and ready for analysis. Below are the key techniques used for cleaning data:

1. Handling Missing Data

Missing values can distort analysis and machine learning models. Common methods to handle them include:

- **Removal:** Delete rows or columns with missing values (only if the missing data is minimal).
 - **Imputation:** Fill in missing values using:
 - Mean, median, or mode (for numerical data)
 - Forward fill or backward fill (for time-series data)
 - Predictive imputation using machine learning models (e.g., KNN imputer)
-

2. Handling Duplicate Data

- Identify duplicate records using pandas:

```
df.duplicated().sum() # Count duplicates
```

```
df.drop_duplicates(inplace=True) # Remove duplicates
```

- Ensure unique identifiers (e.g., IDs) do not have duplicates.
-

3. Handling Outliers

Outliers can skew the results. Methods to handle them:

- **Z-score Method:** Remove values that are too many standard deviations from the mean.
- **IQR Method:** Remove values outside the Interquartile Range (IQR).

```
Q1 = df['column'].quantile(0.25)
```

```
Q3 = df['column'].quantile(0.75)
```

$IQR = Q3 - Q1$

`df = df[(df['column'] >= (Q1 - 1.5 * IQR)) & (df['column'] <= (Q3 + 1.5 * IQR))]`

- **Capping/Winsorization:** Replace extreme values with upper/lower threshold values.
-

4. Standardizing Data Formats

Ensure consistency in data formats:

- Convert dates to a uniform format:

`df['date_column'] = pd.to_datetime(df['date_column'])`

- Standardize text case (e.g., lowercase all entries):

`df['column'] = df['column'].str.lower()`

- Remove special characters in categorical data:

`df['column'] = df['column'].str.replace('[^A-Za-z0-9]+', '', regex=True)`

5. Fixing Data Inconsistencies

- Check for spelling errors or inconsistent labels:

`df['column'].unique()`

- Replace incorrect values:

`df['column'].replace({'wrong_value': 'correct_value'}, inplace=True)`

6. Handling Categorical Data Issues

- Remove leading/trailing spaces in text:

`df['column'] = df['column'].str.strip()`

- Encode categorical values (if required for modeling):

`df = pd.get_dummies(df, columns=['categorical_column'])`

7. Handling Inconsistent Units & Scaling

- Convert units to a common standard (e.g., kg to grams, USD to EUR).
- Normalize or standardize numerical data:

```
from sklearn.preprocessing import MinMaxScaler  
scaler = MinMaxScaler()  
df[['column']] = scaler.fit_transform(df[['column']])
```

8. Handling Structural Errors

- Check column names for typos and standardize them.
- Ensure data types are correct:

```
df['column'] = df['column'].astype('int') # Convert to integer
```

Data Transformation

1.Data Normalization

Normalization scales numeric data to a fixed range, usually [0,1] or [-1,1], ensuring that all features contribute equally to the model.

Min-Max Scaling (0 to 1)

$$X' = \frac{X - X_{\min}}{X_{\max} - X_{\min}}$$

- X is the original value,
- Xmin is the minimum value in the dataset,
- Xmax is the maximum value in the dataset,
- X' is the scaled value.

```
from sklearn.preprocessing import MinMaxScaler  
import numpy as np
```

```
# Sample data
```

```
data = np.array([[10], [20], [30], [40], [50]])
```

```
# Initialize MinMaxScaler
```

```
scaler = MinMaxScaler(feature_range=(0, 1))
```

```
scaled_data = scaler.fit_transform(data)
```

```
print(scaled_data)
```

```
[[0. ]
```

```
 [0.25]
```

```
 [0.5 ]
```

```
 [0.75]
```

```
 [1.  ]]
```

Z-score Normalization (Standardization)

$$X' = (X - \mu) / \sigma$$

Where:

- X = Individual data point
- μ = Mean of the dataset
- σ = Standard deviation of the dataset

```
import numpy as np
```

```
import pandas as pd
```

```
from scipy.stats import zscore
```

```
# Sample dataset
```

```
data = {'Values': [50, 55, 53, 58, 54, 52, 110, 49, 56, 51]}
```

```
df = pd.DataFrame(data)
```

```
# Calculate Z-score
```

```
df['Z-score'] = zscore(df['Values'])
```

```
print(df)
```

```
  Values  Z-score
```

```
0    50 -0.49
```

```
1    55  0.06
```

```
2    53 -0.17
```

```
3    58  0.37
```

```
4    54 -0.06
```

```
5    52 -0.28
```

```
6    110 3.72 # Outlier
7     49 -0.61
8     56 0.17
9     51 -0.39
```

110 has a Z-score of 3.72, meaning it is an outlier.

2. Data Discretization (Binning)

Converts continuous variables into discrete bins.

Equal-Width Binning

Divides data into equal-sized bins.

```
df['binned_column'] = pd.cut(df['column'], bins=5, labels=False)
```

Equal-Frequency Binning

Each bin has approximately the same number of data points.

```
df['binned_column'] = pd.qcut(df['column'], q=5, labels=False)
```

Data Discretization Using Pandas

Let's apply **Equal-Width Binning** and **Equal-Frequency Binning** to a dataset.

Step 1: Create a Sample Dataset

```
import pandas as pd
import numpy as np
```

```
# Generate a sample dataset with continuous numerical values
```

```
np.random.seed(42)
```

```
df = pd.DataFrame({'Age': np.random.randint(18, 80, 15)})
```

```
print(df)
```

Sample Output:

```
   Age
0    22
1    34
2    57
3    45
```

4	29
5	76
6	62
7	19
8	43
9	31
10	50
11	70
12	27
13	38
14	65

Step 2: Apply Equal-Width Binning

Create 3 equal-width bins

```
df['Age_Binned_Equal_Width'] = pd.cut(df['Age'], bins=3, labels=["Young", "Middle-aged", "Senior"])
```

```
print(df[['Age', 'Age_Binned_Equal_Width']])
```

- This method divides the data range into 3 equal-width bins.
- Labels are assigned based on the bin ranges.

Example Output:

	Age	Age_Binned_Equal_Width
0	22	Young
1	34	Young
2	57	Middle-aged
3	45	Middle-aged
4	29	Young
5	76	Senior
6	62	Senior
7	19	Young

8	43	Middle-aged
9	31	Young
10	50	Middle-aged
11	70	Senior
12	27	Young
13	38	Young
14	65	Senior

- Here, **"Young" (18-38)**, **"Middle-aged" (38-58)**, and **"Senior" (58-78)** are the categories.

Step 3: Apply Equal-Frequency Binning

Create 3 bins with approximately equal number of samples

```
df['Age_Binned_Equal_Freq'] = pd.qcut(df['Age'], q=3, labels=["Young", "Middle-aged", "Senior"])
```

```
print(df[['Age', 'Age_Binned_Equal_Freq']])
```

- This method ensures that each bin has (approximately) the same number of data points.

Example Output:

	Age	Age_Binned_Equal_Freq
0	22	Young
1	34	Young
2	57	Middle-aged
3	45	Middle-aged
4	29	Young
5	76	Senior
6	62	Senior
7	19	Young
8	43	Middle-aged
9	31	Young
10	50	Middle-aged
11	70	Senior

12	27	Young
13	38	Middle-aged
14	65	Senior

- The bins here **do not have fixed width** but ensure a nearly equal number of observations.

Data Integration in Data Analytics

Data Integration is the process of combining data from multiple sources into a unified and consistent format to enable efficient analysis and decision-making.

Types of Data Integration Techniques

1. Manual Data Integration

- Involves manually merging datasets from different sources.
- Suitable for small datasets but inefficient for large-scale integration.

2. ETL (Extract, Transform, Load)

- **Extract:** Data is collected from various sources.
- **Transform:** Data is cleaned, formatted, and standardized.
- **Load:** Transformed data is stored in a database or data warehouse.

Example:

```
import pandas as pd

# Load data from different sources
df1 = pd.read_csv("sales_data.csv")
df2 = pd.read_excel("customer_data.xlsx")

# Merge the datasets
df = pd.merge(df1, df2, on="Customer_ID", how="inner")
print(df.head())
```

3. Data Warehousing

- Stores data from different sources in a central repository.
- Used for large-scale analytics (e.g., Amazon Redshift, Google BigQuery).

4. Data Virtualization

- Provides real-time access to multiple data sources without physically moving data.
- Examples: Denodo, IBM Cloud Pak.

5. Schema Integration

- Combines data with different formats or structures into a common schema.
- Used when merging relational databases.

6. Record Linkage (Entity Resolution)

- Identifies and merges duplicate records from different datasets.

Example using Python:

```
import recordlinkage
```

```
indexer = recordlinkage.Index()
```

```
indexer.block('name') # Blocking on 'name' column
```

```
# Compare datasets
```

```
compare = recordlinkage.Compare()
```

```
compare.string('name', 'name', method='jarowinkler', threshold=0.85)
```

```
# Generate matches
```

```
matches = compare.compute(indexer.index(df1, df2), df1, df2)
```

```
print(matches)
```

7. API-Based Integration

- Retrieves real-time data from web services or cloud sources.

Example:

```
import requests
```

```
response = requests.get("https://api.example.com/data")
```

```
data = response.json()
```

```
df = pd.DataFrame(data)
```

Challenges in Data Integration

1. **Schema Mismatch:** Different column names and formats.
2. **Duplicate Data:** Repetitive records across sources.

3. **Data Inconsistency:** Variations in data formats.
4. **Scalability:** Handling large datasets efficiently.

Data Reduction Techniques

(a) Dimensionality Reduction

Reduces the number of features while retaining most of the information.

- **Principal Component Analysis (PCA):**
 - Transforms features into a new set of uncorrelated variables (principal components).
 - Useful for high-dimensional datasets.
- **Linear Discriminant Analysis (LDA):**
 - Similar to PCA but optimizes for class separability in classification problems.
- **Feature Selection Methods:**
 - **Filter Methods** (Correlation, Mutual Information)
 - **Wrapper Methods** (Recursive Feature Elimination)

- **Embedded Methods** (LASSO, Decision Trees)

Principal Component Analysis (PCA) is a dimensionality reduction technique used to transform high-dimensional data into a lower-dimensional space while retaining as much variance as possible. It works by finding a set of new orthogonal axes (principal components) that maximize the variance in the data.

Key Concepts of PCA:

1. **Standardization** – Data is normalized (zero mean, unit variance).
2. **Covariance Matrix Computation** – Finds relationships between features.
3. **Eigen Decomposition** – Extracts eigenvalues and eigenvectors from the covariance matrix.
4. **Principal Components Selection** – Top components (based on eigenvalues) are chosen to reduce dimensions.
5. **Projection** – Data is transformed into the new feature space.

Why Use PCA?

- Reduces **computational cost** and memory usage.
- Helps in **removing multicollinearity** in datasets.
- Improves **visualization** of high-dimensional data.
- Enhances **model performance** by focusing on the most relevant information.

Step-by-Step Working of PCA

1. Standardization of Data (Preprocessing)

Before applying PCA, the dataset should be standardized so that features have **zero mean** and **unit variance** to prevent large-value features from dominating.

Mathematically, for each feature:

$$X_{\text{scaled}} = (X - \mu) / \sigma$$

where:

- X is the original feature value.
- μ is the mean of the feature.
- σ is the standard deviation of the feature.

2. Compute Covariance Matrix

- The covariance matrix captures the relationships between different features. It helps to understand how features vary together.

For two features X_1 and X_2 , covariance is given by:

$$\text{Cov}(X_1, X_2) = \frac{1}{n} \sum_{i=1}^n (X_{1i} - \bar{X}_1)(X_{2i} - \bar{X}_2)$$

3. Compute Eigenvalues and Eigenvectors

- **Eigenvalues:** Measure the variance explained by each principal component.
- **Eigenvectors:** Represent the new axes in transformed space.

$$Cv = \lambda v$$

where:

- λ are the eigenvalues.
- v are the eigenvectors.

4. Sort Eigenvalues and Select Principal Components

- Eigenvalues are sorted in descending order.
- The corresponding top eigenvectors form the **principal components**.

Example: If we choose **top 2 components**, we retain most of the variance.

5. Project Data onto New Feature Space

The original dataset X is projected onto the selected principal components:

$$X_{\text{new}} = X_{\text{scaled}} \times V_k$$

where V_k contains the top k eigenvectors.

(b) Numerosity Reduction

Stores data in a compressed form instead of raw data.

- **Sampling:**

- Selects a subset of data points (e.g., random sampling, stratified sampling).
 - **Regression Models:**
 - Replaces actual data with a mathematical model (e.g., linear regression, polynomial regression).
 - **Histograms:**
 - Summarizes data by grouping values into bins.
 - **Clustering:**
 - Groups similar data points and stores cluster representatives (e.g., K-Means).
-

(c) Data Compression

Encodes data in a compact format.

- **Lossless Compression** (e.g., Huffman Coding, Run-Length Encoding)
 - No information loss, reversible.
 - **Lossy Compression** (e.g., JPEG, MP3)
 - Reduces storage size by removing less important details.
-

Use Cases

- ✓ Speeding up machine learning models
- ✓ Reducing memory and storage requirements
- ✓ Improving model interpretability

Data Cube Aggregation

Data cube aggregation is a multi-dimensional summarization technique used in Online Analytical Processing (OLAP). It enables fast query execution by precomputing aggregations at different levels.

(a) Concept of Data Cubes

A **data cube** organizes data along multiple dimensions, allowing aggregated analysis at various levels.

- Example: **Sales Data Cube** with dimensions:

- **Time** (Year → Quarter → Month → Week)
- **Location** (Country → State → City → Store)
- **Product** (Category → Subcategory → Item)

(b) Aggregation Types

- **Roll-Up (Drill-Up)**: Aggregates data to a higher level (e.g., sum monthly sales to get yearly sales).
- **Drill-Down**: Moves from higher-level aggregation to detailed data (e.g., yearly sales → monthly sales).
- **Slice**: Extracts a subset along one dimension (e.g., sales in 2023).
- **Dice**: Extracts a subset along multiple dimensions (e.g., sales of electronics in Q1 2023).
- **Pivot**: Rotates dimensions for different perspectives.

(c) Benefits

- ✓ Reduces query processing time by precomputing summaries.
- ✓ Optimizes storage by storing aggregated values instead of raw data.
- ✓ Enhances data analysis efficiency in business intelligence.

Data cube aggregation, or roll-up, involves summarizing data across multiple dimensions to create a higher-level view, like aggregating daily sales to monthly or yearly totals. For example, you might aggregate sales data by region, product category, and time period to create a cube showing total sales for each region, product, and year.

Here's a more detailed explanation and example:

What is a Data Cube?

- A data cube is a multidimensional structure used in Online Analytical Processing (OLAP) to store and analyze data.
- It organizes data along multiple dimensions (like time, location, product, etc.) and measures (like sales, revenue, etc.).
- Each cell in the cube represents a specific combination of dimensions and holds a corresponding measure value.

What is Data Cube Aggregation (Roll-up)?

- Data cube aggregation (also called roll-up or drill-up) is the process of summarizing data from a lower level of detail to a higher level.

- This is achieved by combining data across multiple dimensions to create aggregated values.
- For example, you can roll up daily sales data to monthly sales data, or monthly sales data to quarterly sales data.

Example:

Imagine you have sales data for a company, with dimensions of:

- **Product Category:** (e.g., "Electronics", "Clothing", "Food")
- **Region:** (e.g., "North", "South", "East", "West")
- **Time Period:** (e.g., "Daily", "Monthly", "Quarterly", "Yearly")
- **Measure:** (e.g., "Total Sales")

Here's how data cube aggregation (roll-up) would work:

1. **Initial Data:**

You start with detailed sales data, like daily sales for each product category and region.

2. **Roll-up by Time:**

You can roll up the daily sales data to monthly sales data by summing the sales for each product category and region within each month.

3. **Roll-up by Region:**

You can roll up the monthly sales data to regional sales data by summing the sales for each product category within each region.

4. **Roll-up by Product Category:**

You can roll up the regional sales data to product category sales data by summing the sales for each region within each product category.

5. **Final Cube:**

The resulting data cube would contain aggregated sales data, showing total sales for each product category, region, and time period.

Benefits of Data Cube Aggregation:

- **Simplified Analysis:** Aggregated data is easier to analyze and understand than raw data.
- **Improved Performance:** Aggregating data can improve query performance, especially for large datasets.
- **Trend Identification:** Aggregating data helps identify trends and patterns in the data.

- **Data Reduction:** Aggregation helps reduce the volume of data by summarizing it along different dimensions, making it easier to analyze and visualize trends, patterns, and relationships.